


## Problem Set 2: Time and Ordering, Snapshots

1. [5 pts] A clock is reading 12:33:57.0 (hr:min:sec) when it is discovered to be 3 seconds fast. Explain why it is undesirable to set it back to the right time at that point and show (numerically) how it should be adjusted so as to be correct after 6 seconds has elapsed.
2. [5 pts] A client attempts to synchronize with a time server. It records the round-trip times and timestamps returned by the server in the table below. Which of the times in the table below should it use to set its clock and why? What should it set its clock to? What is the estimated accuracy w.r.t to server's clock? If minimum time to send or receive a message from the server is known to be at least 8ms, does your answer change? Why?


Round-trip Time (ms)	Server Clock (hh:mm:ss)
24	12:43:33.456
28	12:43:35.235
22	12:43:36.124

3. [3 pts] If in problem 2 if it is required to sync to within 1ms of the server, is that feasible? Explain your answer.
4. [5 pts] By considering a chain of zero or more messages connecting events  $e$  and  $e'$ , and using induction, show that  $e \rightarrow e'$  implies  $L(e) \leq L(e')$ , where  $L(e)$  is the Lamport timestamp of event  $e$ .
5. [5 pts] Show that  $V_j[i] \leq V_i[i]$  where  $V_i$  is vector clock at process  $i$ .
6. [5 pts] In a similar fashion to 4, show that  $e \rightarrow e'$  implies  $V(e) \leq V(e')$
7. [5 pts] Using result from 5, show that if  $e$  and  $e'$  are concurrent that neither  $V(e) \leq V(e')$  nor  $V(e') \leq V(e)$ . Thus show that if  $V(e) < V(e')$  then  $e \rightarrow e'$ .


I. Mark one pair of events that are concurrent with each other.




II. Mark one pair of events that are concurrent with each other.




iii. Identify an event at process P2 that happens before an event at P1.




IV. Mark all Lamport timestamps on this figure for all events.  
All processes start with zero timestamps.




V. Mark all vector timestamps on this figure for all events.  
All processes start with zero timestamps.



VI. Mark all Lamport timestamps on this figure for all events.  
All processes start with zero timestamps.




VII. Mark all vector timestamps on this figure for all events.  
All processes start with zero timestamps.





- Mark the entire global snapshot collected.



P2 is initiator

- Record local state  $S_2 \xrightarrow{C}$
- Sends out markers
- Turn on recording on channels  $C_0, C_1$

 Marker Message  
 Regular Message  $c$

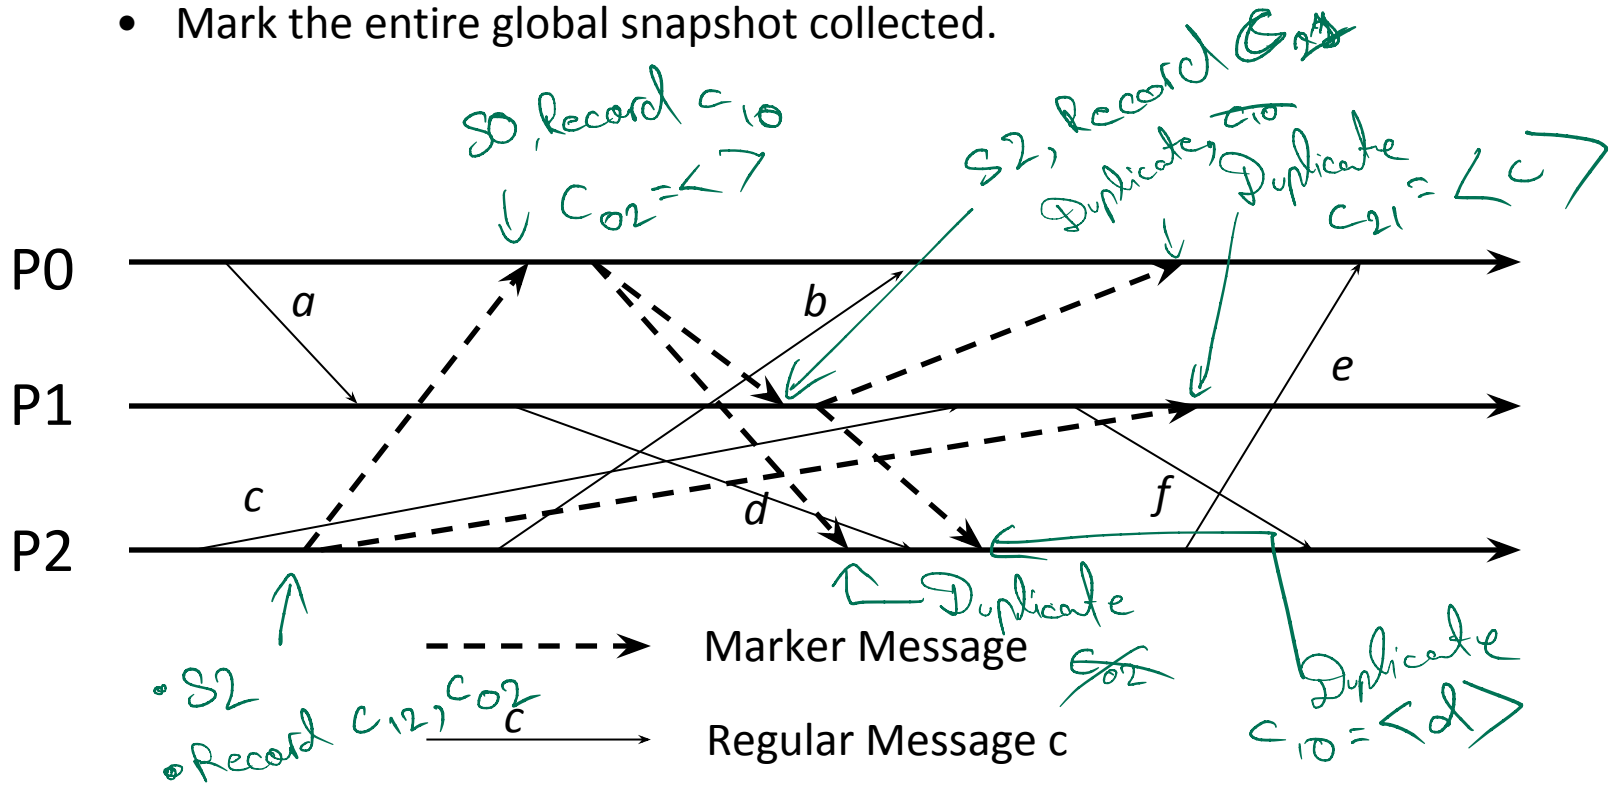
marked

↑ stops recording ~~60~~

$$\begin{aligned} S_{12} C_{01} &= \langle \rangle & S_{21} C_{12} &= \langle \rangle \\ S_{02} C_{10} &= \langle \rangle & S_{12} C_{21} &= \langle \rangle \\ S_{01} C_{01} &= \langle \rangle \\ S_{21} C_{21} &= \langle \rangle \end{aligned}$$

## II. Chandy-Lamport Global Snapshots Algorithm

- Mark the entire global snapshot collected.




$$S_2 = C_{02} = \langle \rangle, C_{12} = \langle \rangle$$

$$S_0 = c_{10} = \langle d \rangle, C_{12} = \langle \rangle$$

$$S_1 = C_{21} = \langle c \rangle$$


III. Mark all Lamport timestamps for application messages on this figure for all events.

All Lamport timestamps start from zero.



IV. Mark all vector timestamps for application messages on this figure for all events.

All vector timestamps start from zeroes.



## Problem Set 2: Time and Ordering, Snapshots

1. A clock is reading 12:33:57.0 (hr:min:sec) when it is discovered to be 3 seconds fast. Explain why it is undesirable to set it back to the right time at that point and show (numerically) how it should be adjusted so as to be correct after 6 seconds has elapsed.

It will be undesirable to set the clock back to the right time as we will have stored time events in the future in our log which will cause the system to malfunction. We will fix this by slowing down our clock. To fix the 3 second drift in 6 seconds, we will slow down our clock by a factor of

$$\frac{1}{2}(T_{\text{error}} + 6\text{seconds} - 12:33:57:0) + 12:33:57:0.$$

2. A client attempts to synchronize with a time server. It records the round-trip times and timestamps returned by the server in the table below. Which of the times in the table below should it use to set its clock and why? What should it set its clock to? What is the estimated accuracy w.r.t to server's clock? If minimum time to send or receive a message from the server is known to be at least 8ms, does your answer change? Why?

Round-trip Time (ms)	Server Clock (hh:mm:ss)
24	12:43:33.456
28	12:43:35.235
22	12:43:36.124

The client should use the time corresponding to the minimum round trip time of 22 seconds.

The error for the case with no network latency will be  $(22-0-0)/2 = 11\text{ms}$  according to Cristian's algorithm. The time set by the algorithm will be  $12:43:36.124 + 0.011 = 12:43:36.135$

The Error will be at most  $(22-8-8)/2\text{ms} = 3\text{ms}$  with Cristian's algorithm. The time set by the algorithm will be  $12:43:36.124 + \{(22 + 8 - 8)/2 \text{ ms}\} = 12:43:36.124 + 0.011 = 12:43:36.135$

While the overall time remains the same, with the server latency being 8 seconds, the error is bounded to 3ms while the error for a network with no latency is at most 11ms. This is because as the time duration for measuring intervals decreases (as we go towards no latency), measuring time gets harder.

3. If in problem 2 if it is required to sync to within 1ms of the server, is that feasible? Explain your answer.

To synchronize a clock within  $\pm 1$  ms it is necessary to obtain a round-trip time of no more than 18 ms if the network latency is 8 ms. ( $RTT - 8 - 8 = 1$ ,  $RTT = 18$ ms) It is possible to obtain round-trip time of 18ms, but it may be improbable that such a time could be found. The file server risks failing to synchronize over a long period, when it could synchronize with a lower accuracy.

4. By considering a chain of zero or more messages connecting events  $e$  and  $e'$ , and using induction, show that  $e \rightarrow e'$  implies  $L(e) \leq L(e')$ , where  $L(e)$  is the Lamport timestamp of event  $e$ .
 

Step 1: The event  $e_1 \rightarrow e_2$  will be true if there is a direct connection between the two nodes and will imply  $L(e_1) \leq L(e_2)$  (Will be true as the event  $e_1$  will send its value (1 in this case) which will be incremented by  $e_2$ )

Step 2: Let  $e_1 \rightarrow e_2 \rightarrow \dots \rightarrow e_{k-2} \rightarrow e_{k-1} \rightarrow e_k$  imply  $L(e_1) \leq L(e_2) \leq \dots \leq L(e_{k-2}) \leq L(e_{k-1}) \leq L(e_k)$  be true for  $k$  causal events

Step 3: To Prove:  $L(e_1) \leq \dots \leq L(e_n)$  will be true for an event  $n = k+1$  that is  $e_n \rightarrow e_{n+1}$  if there is a causal relationship between them

Let there be a chain of events connection  $e$  and  $e'$  such that  $e = e_1$  and  $e' = e_n$ . As these are connected,  $e_1 \rightarrow e_n$  and from the Induction hypothesis,  $L(e_1)$  will be less than  $L(e_n)$ . Hence  $L(e)$  will be less than  $L(e')$ .
5. Show that  $V_j[i] \leq V_i[i]$  where  $V_i$  is vector clock at process  $i$ .
 

Lets us assume that  $V_j[i]$  is strictly greater than  $V_i[i]$ . If this is true then every connection from  $j$  to  $i$  will increase  $i$ 's value in the vector array for  $V_i$ . This will also imply that in the beginning state,  $P_j$  will have detected an event  $e_i$  before  $P_i$  itself but this is not possible as there is no way for  $P_j$  to know the events that occur at Process  $i$  without process  $i$  without process  $i$  recording it. This will also not follow the property of vector clocks that state that "If  $e_j$  occurred before  $p_j$  recorded its state, then  $e_i$  must have occurred before  $p_i$  recorded its state". As this will not be possible, our assumption must be wrong and hence  $V_j[i]$  must always be less than  $V_i[i]$ .
6. In a similar fashion to 4, show that  $e \rightarrow e'$  implies  $V(e) \leq V(e')$ 

Step 1:  $e_1 \rightarrow e_2$  will imply  $V(e_1) \leq V(e_2)$  as  $e_1$  will be connected to  $e_2$  and hence will at least have all the values of the vector smaller than itself from  $e_1$  and will have all the values greater than  $e_1$ 's vector array its existing vector array ( $e_2$ )

Step 2: Let  $e_1 \rightarrow e_k$  imply  $V(e_1) \leq V(e_k)$  be true for all  $k$  (Assuming there is a chain of events leading from  $e_1$  to  $e_k$ )

Step 3: To prove:  $V(e_1) \leq \dots \leq V(e_n)$  where  $n = k+1$ ,  $e_1 = e$  and  $e_n = e'$ .

Let the events  $e_1$  be connected to  $e_n$  by a series of events such  $V(e_1) \leq V(e_k)$  by the induction hypothesis. As  $e_k$  is connected directly to  $e_{k+1}$ ,  $e_k \rightarrow e_{k+1}$  and  $V(e_k) \leq V(e_{k+1})$  is true. This will imply that  $V(e_1) \leq V(e_{k+1})$  and as  $k+1 = n$ ,  $e_1 \rightarrow e_n$  and  $V(e_1) \leq V(e_n)$  (by the Induction Hypothesis).

7. Using result from 5, show that if  $e$  and  $e'$  are concurrent that neither  $V(e) \leq V(e')$  nor  $V(e') \leq V(e)$ . Thus show that if  $V(e) < V(e')$  then  $e \rightarrow e'$ .

Let  $e$  and  $e'$  be concurrent and let  $e$  occur at process  $i$  and  $e'$  occur at process  $j$ . Because these events are concurrent, none of the messages sent from the  $i$ 'th process sent after event  $e$  (inclusive) will have sent its timestamp to process  $j$  by the time  $e'$  occurs at  $p_j$ , and vice versa. We know that the source increments its value in the vector clock by 1 before sending a message and the receiving event increments its value by 1 only after receiving the event (and other entries if their timestamps are greater than the receiving events), it follows that  $V_j[k] < V_i[k]$  and  $V_i[j] < V_j[j]$  and therefore that neither  $V(e) \leq V(e')$  nor  $V(e') \leq V(e)$ .

This can only be possible when the two events  $V(e) < V(e')$  are causal. Hence it is proved that  $e \rightarrow e'$ .

I. Mark one pair of events that are concurrent with each other.






P0


P1

P2

**II. Mark one pair of events that are concurrent with each other.**

P3






P0

P1


P2


P3



IV. Mark all Lamport timestamps on this figure for all events. All processes start with zero timestamps.








P0

P1

P2


P3



**VII. Mark all vector timestamps on this figure for all events. All processes start with zero timestamps.**







# I. Chandy-Lamport Global Snapshots Algorithm•


Mark the entire global snapshot collected.

— — — ➔

————— ➔

$c$

Regular Message  $c$



Marker Message

## II. Chandy-Lamport Global Snapshots Algorithm•


Mark the entire global snapshot collected.

— — — →

—————→

*c*

Regular Message *c*



III. Mark all Lamport timestamps for application messages on this figure for all events.


All Lamport timestamps start from zero.

— — — →

————→

*c*

Regular Message *c*



IV. Mark all vector timestamps for application messages on this figure for all events.



All vector timestamps start from zeroes.

— — — →

————→

*c*

Regular Message *c*



*c* Regular Message *c*